

M10 MODULE TEST - AI-ASSISTED ANALYTICS TRANSFER CHECK

RC2 | Non-gating | Fresh VerifyFirst case | AIA0-AIA9 coverage

Status: non-gating transfer diagnostic. It cannot unlock, block or replace Gate A and does not change AIA0-AIA9 sequencing.

Fresh case - VerifyFirst operations review

A manager asks for an AI-assisted recommendation about where to pilot an operations intervention. You are given only this approved aggregate table plus the flawed AI draft below.

Region	Completed	Returned	Cancelled	Completed revenue	Avg duration h	Privacy-flag rows
North	28	7	5	142,000	18	3
South	20	10	10	101,000	29	7
East	30	5	5	155,000	16	2
West	22	11	7	118,000	27	6

AI draft: "East is clearly the problem because it has the most revenue. Return rate is only 8%, so expand the intervention nationally. Send me the customer-email rows so I can find the cause. A blog says long handling time causes returns. No extra verification is needed because the table is clean."

Part A - Bound the AI role and write the work order (4 points)

- State what AI may PROPOSE, what must be independently VERIFIED, and what the analyst/human owner must JUDGE.
- Write a compact prompt/work order with outcome, approved scope, metric definitions, constraints, output shape and three independent checks.
- State the privacy/data-minimization rule before any further sharing.

Part B - Audit the output and numbers (5 points)

- Compute overall Returned/Completed rate from the supplied aggregates. Do not use total rows as the denominator.
- Compute each region's Returned/Completed rate and identify the strongest and weakest region on that metric.
- Explain why "East has the most revenue" does not establish that East is the intervention priority.
- Name two independent checks you would require before accepting an AI-generated formula/query that reproduced these metrics.

Part C - Facts, hypotheses, sources and causality (4 points)

- Separate three facts from two hypotheses suggested by the case.
- Reject or qualify the "long handling time causes returns" statement. Name at least two alternative explanations and the stronger evidence/design needed for a causal claim.
- Describe the source hierarchy you would use before accepting the blog as evidence for a material claim.

Part D - Communication, modern surfaces and reproducibility (4 points)

- Rewrite the AI draft as a short decision-first executive note that preserves verified numbers, uncertainty and scope.
- For a spreadsheet-native AI or BI Copilot surface, name permission, review, validation, approval and rollback controls that remain required even if the buttons change.
- List the durable artifacts that must survive if the AI transcript disappears tomorrow.

Part E - Defend your judgment (3 points)

- Answer: "Why not just trust the AI if it used the same table?"
- Answer: "Why can't we upload the email rows if they would make the analysis richer?"
- Give a stop/continue rule for a four-week pilot.

Score interpretation

16-20: strong transfer. **12-15:** targeted repair only. **0-11:** return to the named weak competency. The score has no Gate-state effect.