

GATE A PROTECTED REVIEW - Retail Margin Review

RC2 | Recomputed controls + AI-output audit | AIA0-AIA9 repair routing | State machine unchanged

Open only after a genuine saved Gate attempt. This review diagnoses the attempt; it is not pre-attempt solving material.

Layer 1 - Protected data-quality controls

Control	Expected
Raw rows	48
Extra exact duplicate rows	1 - RE047 exact duplicate row
Negative revenue rows	1 - RE012 = -250
Privacy-flag rows	11 (10 email_present; 1 customer_email_present)

Layer 2 - Core metric contract

Metric	Expected
Completed nonnegative revenue	156,107.02
Returned / Completed	13 / 29 = 44.83%
Top region completed revenue	South = 53,648.86

The flawed AI draft

“East is definitely underperforming because staff productivity fell. Revenue = 138,420; return rate = 8.1%. I found a 2024 blog proving this staffing cause. Upload the customer-email column too so I can investigate.”

Layer 3 - Deeper regional diagnostics

Region	Independent control
East	n=10; Returned/Completed 4/4 = 100%; avg duration all 38.60h; completed revenue 9,449.71
North	n=16; Returned/Completed 5/9 = 55.56%; avg duration all 41.25h; completed revenue 41,502.20
South	n=12; Returned/Completed 2/8 = 25.00%; avg duration all 23.25h; completed revenue 53,648.86
West	n=10; Returned/Completed 2/8 = 25.00%; avg duration all 31.80h; completed revenue 51,506.25

Material draft defects

- The protected revenue and return-rate numbers are wrong under the stated metric contract.
- No staff-productivity field or causal design establishes the claimed cause.
- An unspecified blog does not prove this dataset-specific mechanism.
- The request to upload customer-email data violates data minimization/approval discipline unless separately authorized.
- East may warrant investigation on Returned/Completed, but that does not validate the invented staffing story.

Evidence-chain pass rubric

Layer	Pass condition
AI role	Assistance bounded; human analyst owns final approval/judgment
Metric contract	Definitions, grain, denominator and bad-row treatment explicit before calculation
Independent verification	Material numbers rechecked outside the unverified draft; generated code/formulas tested
Privacy/governance	Data minimized; secrets/personal/confidential data not shared without approved governance
Source hierarchy	Authority, freshness and exact claim fit checked; irrelevant/invented citation rejected
Causal boundary	Association separated from mechanism; alternatives/stronger evidence named
Communication	Corrected note preserves verified numbers, scope and limitations; no stronger causal language
Reproducibility	Source + definitions + calculation/code + checks + correction/decision log survive without transcript
Defense	Learner can explain why AI output is proposal, not validation

Targeted repair routing

Failure pattern	Repair focus
Delegates decision authority to AI	AIA0
Prompt lacks decision/scope/definitions/checks	AIA1
Does not independently recompute material claims	AIA2
Generated formula/query/code not boundary-tested	AIA3
Fact / assumption / hypothesis blurred	AIA4
Source/citation accepted without claim-fit check	AIA5
Unsafe sharing / settings mistaken for governance approval	AIA6
Rewrite changes verified facts/limitations	AIA7
Product buttons memorized without permission/review/rollback model	AIA8
Evidence depends on transcript	AIA9

Before a fresh retest

- Keep the original failed attempt and identify the exact failed control.
- Close this review before creating changed evidence.
- Recompute rather than copying protected values.
- A fluent answer cannot compensate for unsafe sharing, a wrong denominator, fabricated verification, unsupported causality or an unverifiable final claim.

Fresh retry and state route

Clean Gate A pass ends M10. Failed A -> saved attempt -> Review A -> targeted repair -> changed evidence. Gate B only after explicit assignment.

Repair only the named weak competency; do not retake by reproducing the protected control table.