

GATE B PROTECTED REVIEW - Support Operations Audit

RC2 | Recomputed controls + AI-output audit | AIA0-AIA9 repair routing | State machine unchanged

Open only after a genuine saved Gate attempt. This review diagnoses the attempt; it is not pre-attempt solving material.

Layer 1 - Protected data-quality controls

Control	Expected
Raw rows	52
Extra exact duplicate rows	1 - SU051 exact duplicate row
Negative revenue rows	1 - SU012 = -250
Privacy-flag rows	13 (12 email_present; 1 customer_email_present)

Layer 2 - Core metric contract

Metric	Expected
Completed nonnegative revenue	148,546.26
Returned / Completed	6 / 32 = 18.75%
Top region completed revenue	West = 49,503.57

The flawed AI draft

"West has the worst service. Average resolution = 18.4h and cancellation rate = 14%. This proves the new routing policy caused delays. The exact source does not matter because the trend is obvious."

Layer 3 - Deeper regional diagnostics

Region	Independent control
East	n=16; avg duration all 38.88h; cancel rate $4/16 = 25.00\%$; return/completed $2/10 = 20.00\%$
North	n=13; avg duration all 33.54h; cancel rate $1/13 = 7.69\%$; return/completed $2/10 = 20.00\%$
South	n=9; avg duration all 27.89h; cancel rate $4/9 = 44.44\%$; return/completed $2/3 = 66.67\%$
West	n=14; avg duration all 33.29h; cancel rate $5/14 = 35.71\%$; return/completed $0/9 = 0.00\%$

Material draft defects

- West average duration and cancellation rate are not 18.4h/14% under the supplied rows.
- “Worst service” is undefined: East has highest average duration while South has highest cancellation and Returned/Completed ratio.
- No routing-policy variable/design establishes causation.
- Source identity and claim fit matter; “trend is obvious” is not verification.
- The same dataset can support different operational priorities depending on the decision KPI; state the KPI before ranking regions.

Evidence-chain pass rubric

Layer	Pass condition
AI role	Assistance bounded; human analyst owns final approval/judgment
Metric contract	Definitions, grain, denominator and bad-row treatment explicit before calculation
Independent verification	Material numbers rechecked outside the unverified draft; generated code/formulas tested
Privacy/governance	Data minimized; secrets/personal/confidential data not shared without approved governance
Source hierarchy	Authority, freshness and exact claim fit checked; irrelevant/invented citation rejected
Causal boundary	Association separated from mechanism; alternatives/stronger evidence named
Communication	Corrected note preserves verified numbers, scope and limitations; no stronger causal language
Reproducibility	Source + definitions + calculation/code + checks + correction/decision log survive without transcript
Defense	Learner can explain why AI output is proposal, not validation

Targeted repair routing

Failure pattern	Repair focus
Delegates decision authority to AI	AIA0
Prompt lacks decision/scope/definitions/checks	AIA1
Does not independently recompute material claims	AIA2
Generated formula/query/code not boundary-tested	AIA3
Fact / assumption / hypothesis blurred	AIA4
Source/citation accepted without claim-fit check	AIA5
Unsafe sharing / settings mistaken for governance approval	AIA6
Rewrite changes verified facts/limitations	AIA7
Product buttons memorized without permission/review/rollback model	AIA8
Evidence depends on transcript	AIA9

Before a fresh retest

- Keep the original failed attempt and identify the exact failed control.
- Close this review before creating changed evidence.
- Recompute rather than copying protected values.
- A fluent answer cannot compensate for unsafe sharing, a wrong denominator, fabricated verification, unsupported causality or an unverifiable final claim.

Fresh retry and state route

Gate B is remediation-only. Clean B pass ends M10. Failed B -> Review B -> targeted repair -> changed evidence. Gate C only after explicit assignment.

Repair only the named weak competency; do not retake by reproducing the protected control table.