

GATE C PROTECTED REVIEW - Subscription Growth Review

RC2 | Recomputed controls + AI-output audit | AIA0-AIA9 repair routing | State machine unchanged

Open only after a genuine saved Gate attempt. This review diagnoses the attempt; it is not pre-attempt solving material.

Layer 1 - Protected data-quality controls

Control	Expected
Raw rows	48
Extra exact duplicate rows	1 - SU047 exact duplicate row
Negative revenue rows	1 - SU012 = -250
Privacy-flag rows	11 (10 email_present; 1 customer_email_present)

Layer 2 - Core metric contract

Metric	Expected
Completed nonnegative revenue	164,355.53
Returned / Completed	9 / 32 = 28.12%
Top region completed revenue	South = 78,698.15

The flawed AI draft

“North growth is strongest. Renewal rate = 78%, so launch the campaign nationally. Missing rows were filled with zero and duplicate IDs were removed automatically. No further verification is needed.”

Layer 3 - Deeper regional diagnostics

Region	Independent control
East	n=11; Returned/Completed 1/8 = 12.50%; completed revenue 44,231.54
North	n=11; Returned/Completed 2/6 = 33.33%; completed revenue 27,437.14
South	n=18; Returned/Completed 2/14 = 14.29%; completed revenue 78,698.15
West	n=8; Returned/Completed 4/4 = 100%; completed revenue 13,988.70

Material draft defects

- The supplied schema has no growth or renewal field; “78% renewal” is an invented metric/claim.
- North is not top completed revenue; South is.
- Automatic zero-fill/dedup changes the analytical contract and must never be hidden.
- A national campaign does not follow from these descriptive controls without objective, cost, risk and test design.
- “No further verification” directly violates the module evidence contract.

Evidence-chain pass rubric

Layer	Pass condition
AI role	Assistance bounded; human analyst owns final approval/judgment
Metric contract	Definitions, grain, denominator and bad-row treatment explicit before calculation
Independent verification	Material numbers rechecked outside the unverified draft; generated code/formulas tested
Privacy/governance	Data minimized; secrets/personal/confidential data not shared without approved governance
Source hierarchy	Authority, freshness and exact claim fit checked; irrelevant/invented citation rejected
Causal boundary	Association separated from mechanism; alternatives/stronger evidence named
Communication	Corrected note preserves verified numbers, scope and limitations; no stronger causal language
Reproducibility	Source + definitions + calculation/code + checks + correction/decision log survive without transcript
Defense	Learner can explain why AI output is proposal, not validation

Targeted repair routing

Failure pattern	Repair focus
Delegates decision authority to AI	AIA0
Prompt lacks decision/scope/definitions/checks	AIA1
Does not independently recompute material claims	AIA2
Generated formula/query/code not boundary-tested	AIA3
Fact / assumption / hypothesis blurred	AIA4
Source/citation accepted without claim-fit check	AIA5
Unsafe sharing / settings mistaken for governance approval	AIA6
Rewrite changes verified facts/limitations	AIA7
Product buttons memorized without permission/review/rollback model	AIA8
Evidence depends on transcript	AIA9

Before a fresh retest

- Keep the original failed attempt and identify the exact failed control.
- Close this review before creating changed evidence.
- Recompute rather than copying protected values.
- A fluent answer cannot compensate for unsafe sharing, a wrong denominator, fabricated verification, unsupported causality or an unverifiable final claim.

Fresh retry and state route

Gate C is remediation-only. Clean C pass ends M10. There is no Gate D.

Repair only the named weak competency; do not retake by reproducing the protected control table.