

GATE B PROTECTED REVIEW - StreamFit Renewals

RC2 | Recomputed quality/metrics | Evidence-chain diagnosis | State machine unchanged

Open only after a genuine saved Gate attempt. Diagnose the first attempt; never use this as a pre-attempt solution.

Layer 1 - Data-quality controls

Control	Expected
Valid rows	238
Rejected rows	3
Rejected SF0071	Missing Plan
Rejected SF0121	Invalid negative Monthly_Revenue
Rejected duplicate	SF0031 duplicate Subscriber_ID

Layer 2 - Core protected metrics

Metric / slice	Expected
Overall renewal rate	66.8067%
Overall avg support contacts	1.79
Monthly revenue in valid set	215,362
Revenue at risk from non-renewed	63,921
Basic	n=116; renewal 58.6207%; support 1.67; revenue 69,484; risk 28,752
Plus	n=82; renewal 70.7317%; support 1.88; revenue 81,918; risk 23,976
Pro	n=40; renewal 82.5000%; support 1.98; revenue 63,960; risk 11,193

Layer 3 - Deeper diagnostic controls

Slice	Expected
Partner channel	renewal 62.32%; revenue at risk 24,374
Direct channel	renewal 69.37%; revenue at risk 26,966
Referral channel	renewal 67.24%; revenue at risk 12,581
West / North renewal	60.61% / 60.00%
South / East renewal	75.38% / 71.15%

What a strong finding should say

Basic has the weakest plan renewal rate and the largest plan-level revenue at risk, while West/North are weak regions. Prioritization should reconcile rate, volume/revenue exposure and intervention ownership rather than picking the smallest rate in isolation.

Causal boundary

Support contacts correlate with customer difficulty/tenure/plan complexity and may be consequence as well as cause. Do not claim support contacts caused churn without stronger evidence.

Evidence-chain pass rubric

Layer	Pass condition
Decision frame	Names the decision, analytical question, scope/comparison and Definition of Done
KPI contract	Numerator/denominator or formula, exclusions, grain and owner are explicit
Quality	Rejected rows are visible; no silent cleaning
Diagnosis	Prioritized comparisons; weighted/reconciled where required; one independent check
Findings	Three findings each include comparison, magnitude, scope and implication
Causal boundary	Observation/association separated from mechanism; alternatives named
Recommendation	Owner + action + timing + success measure + risk/dependency + stop/continue rule
Communication	Decision-first memo and 3-slide message-headline story
Defense	Answers denominator, causal and missing-data challenges without changing definitions

Targeted repair routing

Failure pattern	Repair focus
Decision unclear / analysis starts with charts	IBA1
Metric/grain/denominator drift	IBA2
Too many low-value cuts / no priority	IBA3-IBA4
Causal overclaim	IBA5
Finding is only a number or chart label	IBA6
Generic recommendation	IBA7
Memo buries decision	IBA8
Decorative or topic-title slides	IBA9
Defense changes definitions or hides uncertainty	IBA10
Broken end-to-end traceability	IBA11

Before a fresh retest

- Keep the original failed attempt.
- Write the exact failed competency and repair evidence.
- Close this review before changed evidence.
- A polished story cannot compensate for wrong metrics or unsupported causal claims.

Fresh retry and state route

Gate B is remediation-only. Clean B pass ends M09. Failed B -> Review B -> targeted repair -> changed evidence. Gate C only after explicit assignment.

Repair only the named weak competency; do not retake by copying the same control table.