

Gate B review - Delivery Performance

1. Unlock condition

Open only after the genuine Gate notebook attempt and evidence are saved. This review provides acceptance controls and repair direction, not final code.

2. Strong attempt architecture

Frame metrics -> preserve/load raw -> detect issues -> create justified analysis set -> validate dimension key -> merge with validate=/indicator -> grouped/time analysis -> purposeful visuals -> independent reconciliation -> findings/recommendation/limitations -> Restart + Run All.

3. Critical failures

Silent deletion/correction, disabled merge validation, hidden unmatched rows, invalid dates/numbers treated as valid, wrong rate denominator, multiplied money, fabricated evidence, unsupported recommendation or notebook hidden state can block a pass.

4. State machine

B is remediation-only; clean B pass ends M07; C only after failed B + repair + explicit assignment.

Protected objective controls

Control	Expected
Raw Rows	88
Duplicate Id Rows	2
Invalid Date Rows	2
Invalid Cost Rows	1
Analysis Rows	83
Unmatched Partner Rows	1
On Time Rate	0.3133
Avg Late Hours	6.51
Avg Cost Per Km	10.19
Top Cost Region	East

Diagnostic route and changed retry

Load/quality	PY2-PY3: shape/types/duplicates/invalid dates/numbers/row-rejection logic.
Grouped metrics	PY4: result grain, denominator, aggregation and reconciliation.
Merge	PY5: cleaned dimension-key uniqueness, validate=, indicator=True, unmatched rows.
Time/reshape/visual	PY6-PY7: date arithmetic, chronological grouping, summary feeding chart.
EDA/communication	PY8: evidence vs causation, findings/recommendation/limitations.
Reproducibility	PY9-PY10: relative paths, outputs, clean Restart + Run All, smallest useful workflow.
5. Fresh retry Close this review. Repair only the named competency, then complete changed evidence in a fresh context. Keep the original attempt. Follow the protected Gate state machine rather than automatically opening the next Gate.	